

Capstone Project - Milestone Report

Fernanda H. P. Gomes

October 31, 2018

Summary

This report provides a short overview of the data to be used for the the Data Science Specialization Capstone project along with a description of plans for the word prediction algorithm.

1- Data loading

Below we will load into R the three files we will use in the prediction algorithm: blogs, news and twitter.

```
#Selecting the folder
setwd("~/Desktop/ Data Science/Coursera Capstone/data")

#loading the necessary packages
library(stringi)
library(tm)
```

```
## Loading required package: NLP
```

```
library(RWeka)
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(ggplot2)
```

```
##
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:NLP':  
##  
##   annotate
```

```
library(quanteda)
```

```
## Package version: 1.3.14
```

```
## Parallel computing: 2 of 4 threads used.
```

```
## See https://quanteda.io for tutorials and examples.
```

```
##  
## Attaching package: 'quanteda'
```

```
## The following objects are masked from 'package:tm':  
##  
##   as.DocumentTermMatrix, stopwords
```

```
## The following object is masked from 'package:utils':  
##  
##   View
```

```
#let's check how large is the data
```

```
cat("en_US.blogs.txt: " , file.info("./final/en_US/en_US.blogs.txt")$size / (1024*1024) , "mb")
```

```
## en_US.blogs.txt: 200.4242 mb
```

```
cat("en_US.news.txt: " , file.info("./final/en_US/en_US.news.txt")$size / (1024*1024) , "mb")
```

```
## en_US.news.txt: 196.2775 mb
```

```
cat("en_US.twitter.txt: " , file.info("./final/en_US/en_US.twitter.txt")$size / (1024*1024) , "mb")
```

```
## en_US.twitter.txt: 159.3641 mb
```

```
#importing the data
blogs<-file("./final/en_US/en_US.blogs.txt","r")
blogs_lines<-readLines(blogs, encoding = "UTF-8",skipNul = TRUE)

news<-file("./final/en_US/en_US.news.txt","r")
news_lines<-readLines(news, encoding = "UTF-8",skipNul = TRUE)

twitter<-file("./final/en_US/en_US.twitter.txt","r")
twitter_lines<-readLines(twitter, encoding = "UTF-8",skipNul = TRUE)
```

2- Data Preparation In order to plan the next steps we will start by understanding the available data and preparing that for some exploratory analysis

```
#first let's get a better understanding of the data that we have
DataStats <- rbind(stri_stats_general(news_lines), stri_stats_general(blogs_lines), s
stri_stats_general(twitter_lines))
DataStats <- as.data.frame(DataStats)
row.names(DataStats) <- c("news", "blogs", "twitter")
DataStats
```

##	Lines	LinesNEmpty	Chars	CharsNWhite
## news	1010242	1010242	203223154	169860866
## blogs	899288	899288	206824382	170389539
## twitter	2360148	2360148	162096241	134082806

```
# as we have a lot of data, let's make a sample to help with exploratory analysis
set.seed(100)
blogs.sample <- sample(blogs_lines, length(blogs_lines) * 0.03)
news.sample <- sample(news_lines, length(news_lines) * 0.03)
lines_sample <- sample(twitter_lines, length(twitter_lines) * 0.03)

corpus.blog <- corpus(blogs.sample) #creating corpus
corpus.news <- corpus(news.sample) #creating corpus
corpus.lines <- corpus(lines_sample) #creating corpus

corpus <- corpus.blog + corpus.news + corpus.lines

summary(corpus)
```

```
## Corpus consisting of 128089 documents, showing 100 documents:
```

```
##
```

##	Text	Types	Tokens	Sentences
##	text1	60	86	7
##	text2	7	7	1
##	text3	7	7	1
##	text4	5	5	2
##	text5	4	4	1
##	text6	29	31	3
##	text7	21	24	2
##	text8	99	168	7
##	text9	39	52	3
##	text10	25	26	1
##	text11	46	64	5
##	text12	12	13	1
##	text13	77	100	4
##	text14	55	82	2
##	text15	142	229	13
##	text16	5	5	1
##	text17	45	56	2
##	text18	69	102	6
##	text19	4	4	1
##	text20	11	12	1
##	text21	122	198	7
##	text22	75	103	4
##	text23	47	72	3
##	text24	27	30	2
##	text25	19	21	1
##	text26	36	43	2
##	text27	17	17	1
##	text28	36	44	4
##	text29	7	7	1
##	text30	14	16	2
##	text31	91	143	8
##	text32	22	27	2
##	text33	6	6	1
##	text34	52	75	5
##	text35	9	9	1
##	text36	9	10	3
##	text37	8	8	1
##	text38	51	67	3
##	text39	63	83	3
##	text40	30	35	4
##	text41	16	18	1
##	text42	2	2	1
##	text43	33	46	4
##	text44	19	20	1
##	text45	12	13	1
##	text46	91	130	9
##	text47	7	7	1
##	text48	78	103	4
##	text49	24	26	2
##	text50	52	65	2
##	text51	45	51	4
##	text52	43	72	4
##	text53	50	79	1
##	text54	92	128	4

##	text55	10	10	1
##	text56	16	17	3
##	text57	13	13	1
##	text58	12	13	1
##	text59	4	4	1
##	text60	13	14	1
##	text61	41	52	3
##	text62	82	102	5
##	text63	6	6	1
##	text64	11	11	1
##	text65	10	11	2
##	text66	7	7	1
##	text67	51	67	5
##	text68	40	42	1
##	text69	1	1	1
##	text70	63	113	4
##	text71	16	18	2
##	text72	8	8	1
##	text73	4	4	1
##	text74	5	5	1
##	text75	17	21	1
##	text76	6	6	1
##	text77	7	7	1
##	text78	85	116	6
##	text79	3	3	1
##	text80	5	5	1
##	text81	39	51	1
##	text82	58	80	2
##	text83	29	38	2
##	text84	34	38	2
##	text85	10	10	1
##	text86	44	63	5
##	text87	50	57	2
##	text88	112	169	7
##	text89	30	31	1
##	text90	82	126	4
##	text91	39	52	1
##	text92	11	11	1
##	text93	89	156	5
##	text94	35	43	2
##	text95	35	47	2
##	text96	14	14	2
##	text97	3	3	2
##	text98	68	104	4
##	text99	4	4	1
##	text100	33	43	2
##				
##	Source: Combination of corpuses corpus.blog + corpus.news and corpus.lines			
##	Created: Mon Dec 10 23:02:01 2018			
##	Notes:			

```
#creating tokens and cleaning the data

unigram <- tokens (corpus, remove_numbers = TRUE, remove_punct = TRUE, remove_separators = TRUE , ngrams =1 )

bigram <- tokens (corpus, remove_numbers = TRUE, remove_punct = TRUE, remove_separators = TRUE , ngrams =2 )

trigram <- tokens (corpus, remove_numbers = TRUE, remove_punct = TRUE, remove_separators = TRUE , ngrams =3 )

#creating the matrix
dfm.uni <- dfm (unigram, remove = stopwords("english"))
dfm.uni <- dfm_sort(dfm.uni) [,1:40]

dfm.bi <- dfm (bigram, remove = stopwords("english"))
dfm.bi <- dfm_sort(dfm.bi) [,1:40]

dfm.tri <- dfm (trigram, remove = stopwords("english"))
dfm.tri <- dfm_sort(dfm.tri) [,1:40]
```


3- Exploratory Analysis

```
#Ploting top 40 by frequency

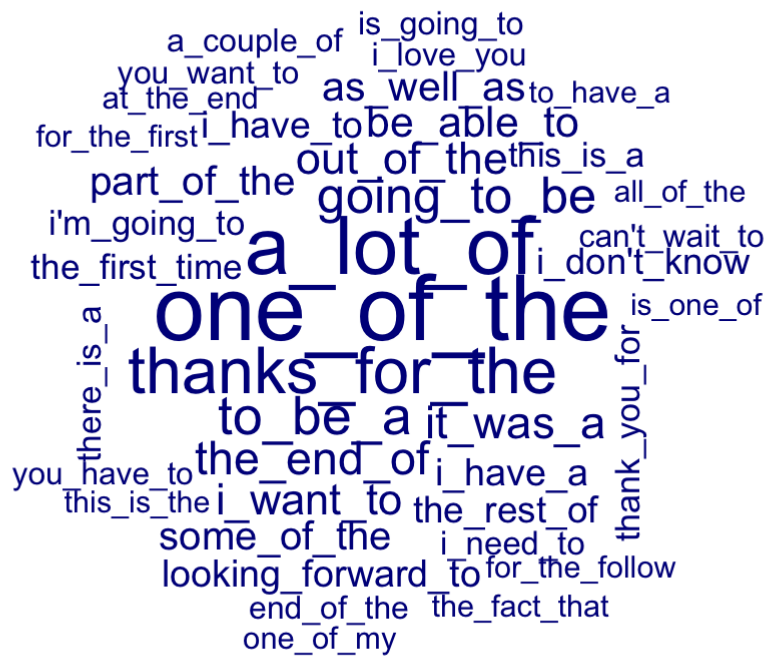
textplot_wordcloud(dfm.uni)
```



```
textplot_wordcloud(dfm.bi)
```



```
textplot_wordcloud(dfm.tri)
```



4 - Prediction algorithm

The goal is to create: - a prediction model based on the n-gram models build - a Shiny app (interface with the user): as the user enter words in a single textbox, the algorithm will be triggered on providing a list of suggested word that the user can select.

Challenge so far: - the dataset is really large, demanding a strategy on how better use the data without killing performance

Next steps: - work on further data cleanse - better define sample selection - build prediction algorithm - test and train datasets for a later prediction model